



D u k e S y s t e m s

CPS 310 / ECE 353


Programs and Processes

Jeff Chase

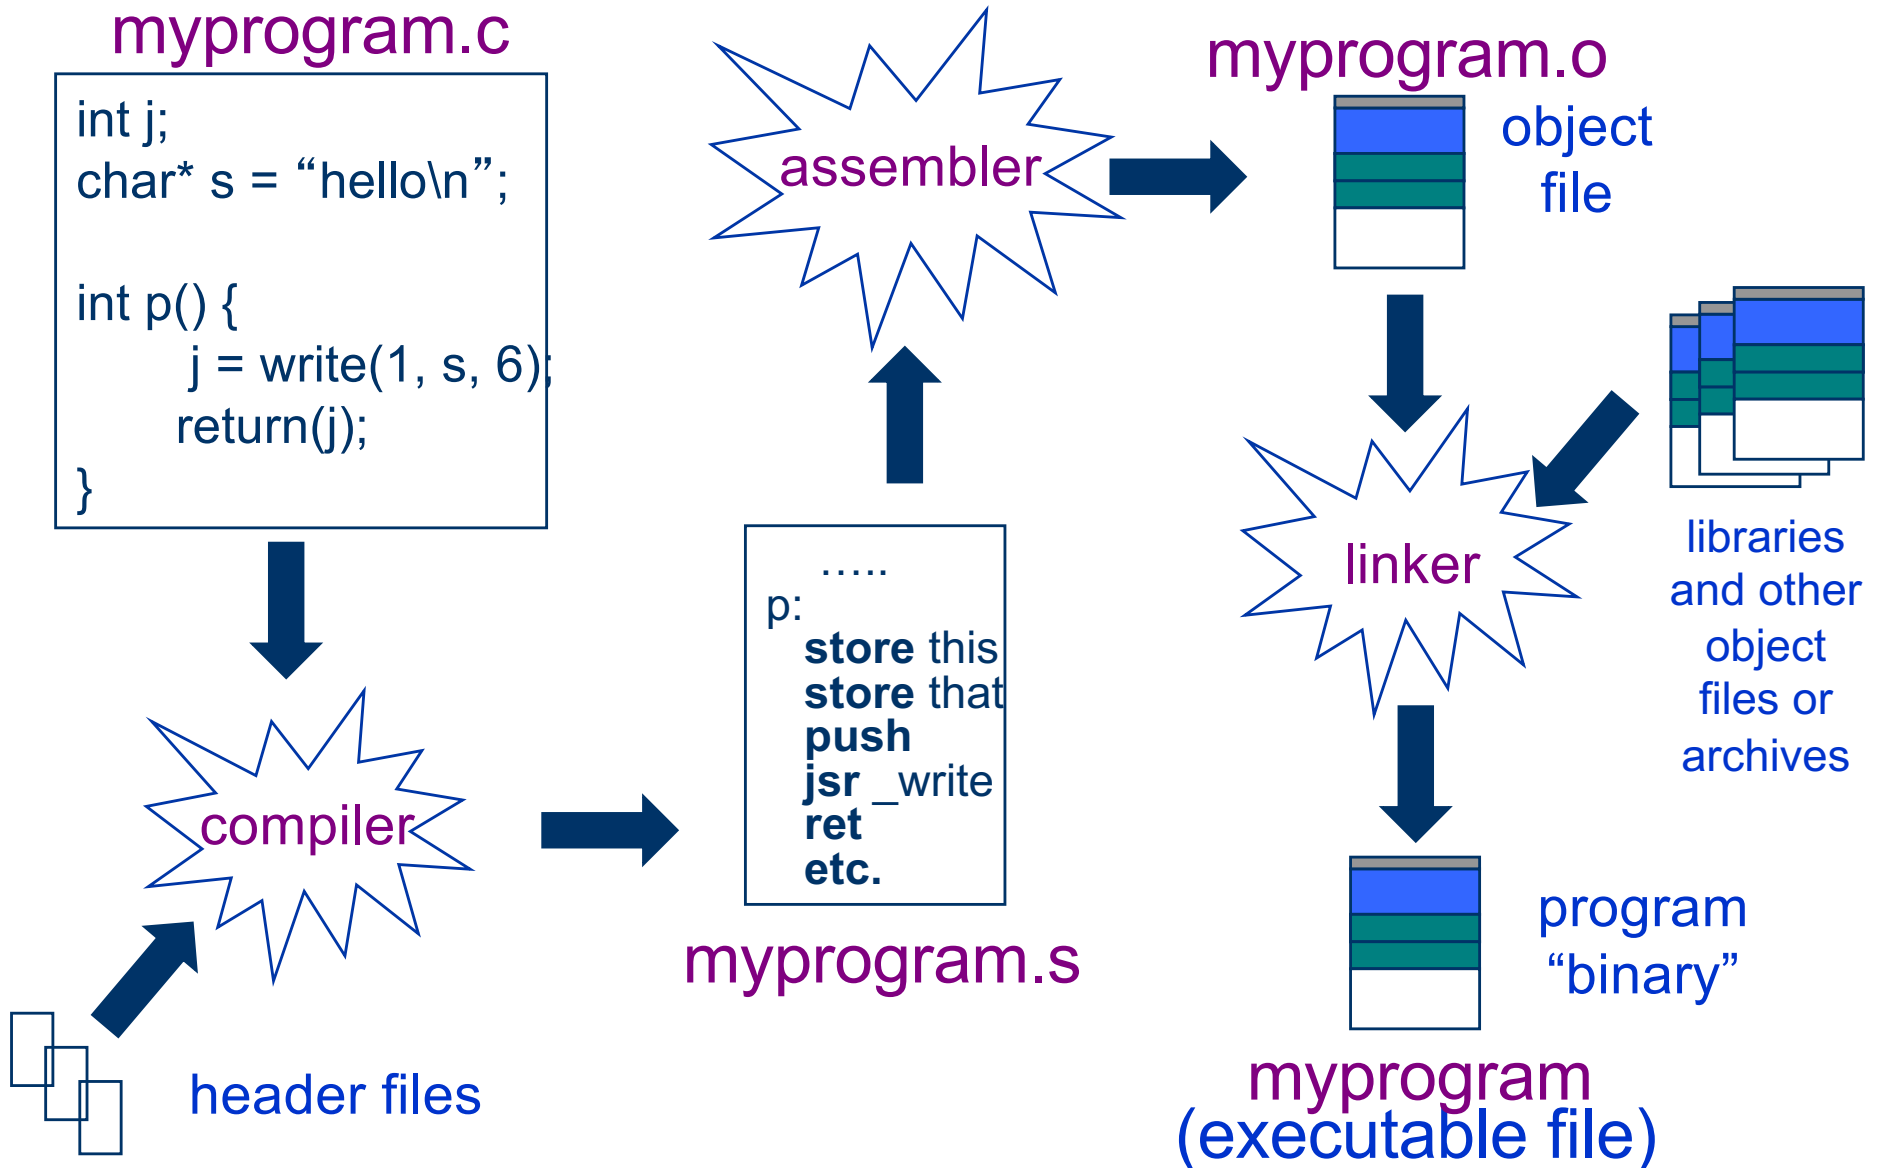
Duke University

What is a program?

- Is it a stupid question?
- A program is data that describes a process. ;-)
- For our purposes today, the data is stored in a file, in a format that the operating system can **execute**.
 - An **executable** or “**binary**” file
- Let's suppose that it is **statically linked** (self-contained).



Birth of a program (C/Ux)

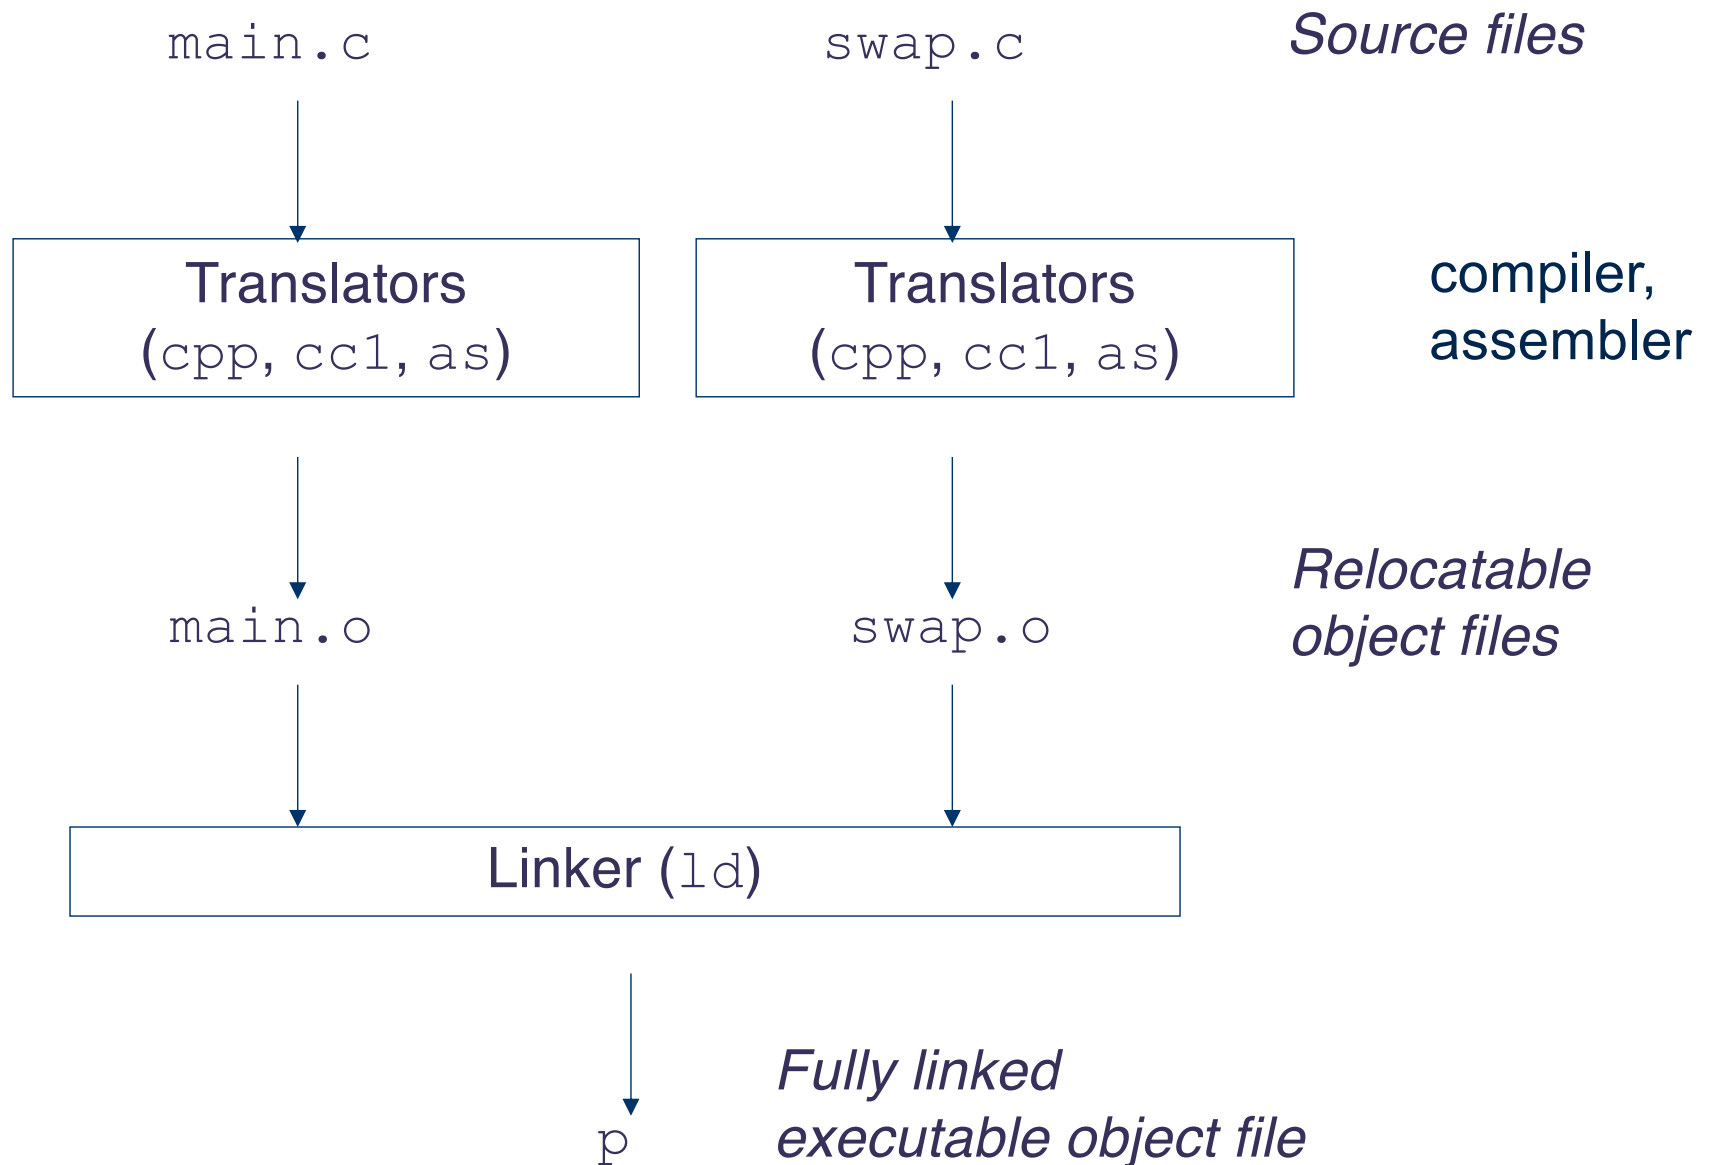


Role of the linker in the build toolchain

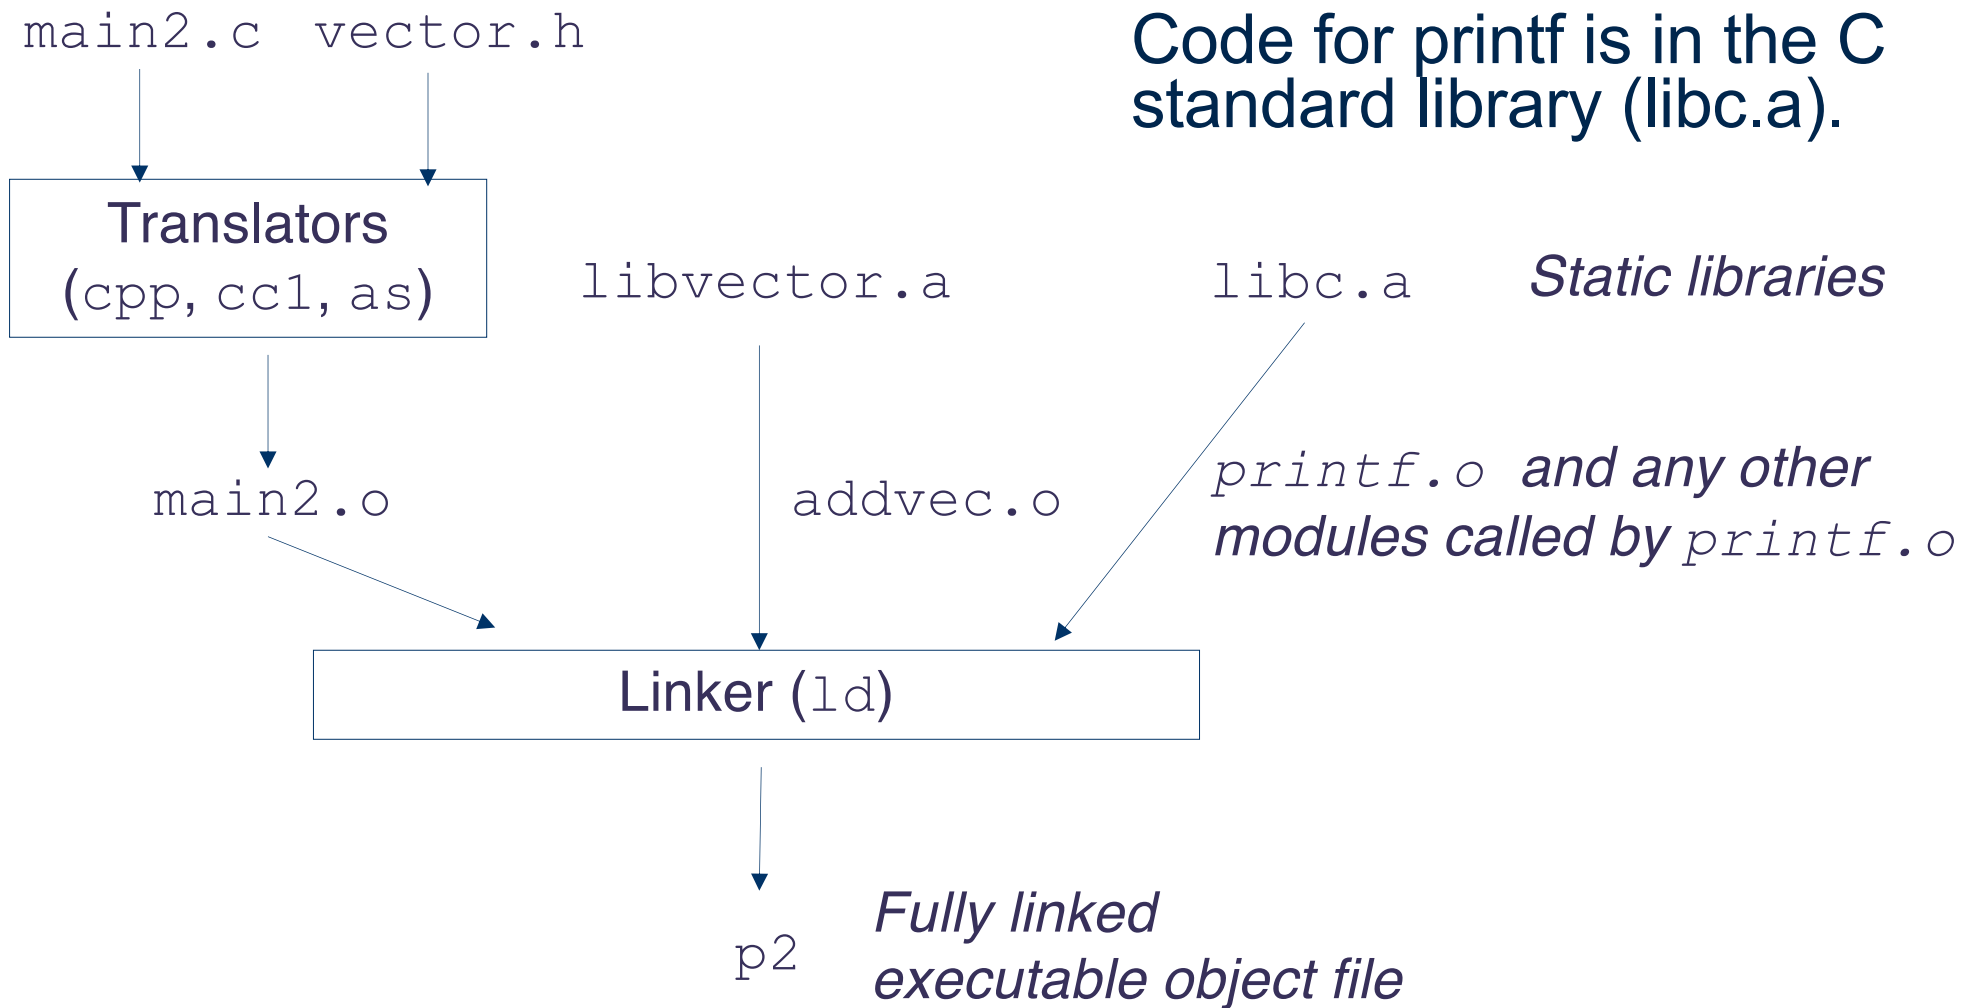
What's this all about?

- Real programs are written in pieces by different people.
- Their code is spread across multiple files.
- They include “off the shelf” code used by other programs.
 - **Libraries**, e.g., the C **standard library** (libc, stdlib).
- We can compile the pieces separately, at different times.
- A piece might reference symbols defined in another.
- The linker combines the pieces into a program.
- C compiler also runs linker for you (unless `-c`), with stdlib.

Separate compilation and the linker



Static linking with libraries (archives)



Where to start? main()

```
chase$ cc -c empty.c
chase$ file empty.o
empty.o: Mach-O 64-bit object x86_64
chase$ nm empty.o
000000000000000000 T _p
chase$ cc empty.c
Undefined symbols for architecture x86_64:
  "_main", referenced from:
      implicit entry/start for main executable
ld: symbol(s) not found...
clang: error: linker failed with exit code 1
chase$
```

empty.c

```
void
p() {}
```

This program compiles OK. It defines the procedure p: a symbol of type T for “text” (code). But it is not a complete program: the **linker doesn't know where to start!** Programs start in main(), but it is missing.

Where to go? Calling out

```
chase$ cc -o nop nop.c
Undefined symbols...:
  "_p", referenced from:
      _main in nop-bbc335.o
ld: symbol(s) not found...
clang: error: linker command failed...
chase$ cc -o nop nop.c empty.o
chase$ ./nop
chase$
```

empty.c

```
void
p() {}
```

nop.c

```
void p();
int main()
    {p();}
```

This program doesn't do anything, but it is a complete program! It has a `main()`. But it also needs `empty.o` to define the symbol `p`, because `main()` calls `p()`. Given both files, the linker is happy. **And we want the linker to be happy.**

What's a library?

- A **library** is just a collection of object files bundled together into a single archive file...
- ...with an index for fast linking.
- **ar** command utility creates/modifies archives.
- No mystery; no magic!

```
chase$ ar cr empty.a empty.o
chase$ file empty.a
empty.a: current ar archive random library
chase$ cc -o nop nop.c empty.a
chase$ ./nop
chase$
```

What's in an object file or executable?

e.g., an ELF or Mach-O file

Header “magic number” indicates type of file/image.


Section table an array of (**offset, len, startVA**)

sections

Metadata used by tools.

Options:

- Remove after final link step and strip (compact).
- Or: add more symbol info for debugger.



program instructions
p

immutable data (constants)
“hello\n”

writable global/static data
j, s, sbuf


j, s, p, sbuf

```
int j = 327;
char* s = “hello\n”;
char sbuf[512];

int p() {
    int k = 0;
    j = write(1, s, 6);
    return(j);
}
```

Sections are byte offset ranges within the file.

ELF: read it like a book



An executable/linkable file (e.g., ELF) is “like a book”.

- **Sections.** They are “like chapters”, or...sections.
 - Each section has content of a given purpose/use, needed to link/run the program.
 - E.g., code, initial values of global data, constants.
- **Header.** Cover page and table of contents (ToC).
- **Symbol table.** Detailed ToC: all defined symbols.
- **Relocation records.** Index: lists references to symbols.

Inside ELF



ANGE ALBERTINI
<http://www.corkami.com>

EXECUTABLE AND LINKABLE FORMAT

```
me@nux:~$ ./mini
me@nux:~$ echo $?
42
```

```

 0 1 2 3 4 5 6 7 8 9 A B C D E F
00: 7F .E .L .F 01 01 01
10: 02 00 03 00 01 00 00 00 60 00 00 08 40 00 00 00
20:                34 00 20 00 01 00

40: 01 00 00 00 00 00 00 00 00 00 00 08 00 00 00 08
50: 70 00 00 00 70 00 00 00 05 00 00 00

60: BB 2A 00 00 00 B8 01 00 00 00 CD 80
```

MINI

Linux uses ELF.
MacOS uses Mach-O.

ELF HEADER

IDENTIFY AS AN ELF TYPE
SPECIFY THE ARCHITECTURE

FIELDS	VALUES
e_ident	
EI_MAG	0x7F, "ELF"
EI_CLASS, EI_DATA	1ELFCLASS32, 1ELFDATA2LSB
EI_VERSION	1EV_CURRENT
e_type	2ET_EXEC
e_machine	3EM_386
e_version	1EV_CURRENT
e_entry	0x8000060
e_phoff	0x0000040
e_ehsize	0x0034
e_phentsize	0x0020
e_phnum	0001

PROGRAM HEADER TABLE

EXECUTION INFORMATION

p_type	1PT_LOAD
p_offset	0
p_vaddr	0x8000000
p_paddr	0x8000000
p_filesz	0x0000070
p_memsz	0x0000070
p_flags	5PF_R PF_X

CODE


X86 ASSEMBLY

```
mov ebx, 42
mov eax, SC_EXIT1
int 80h
```

EQUIVALENT C CODE

```
return 42;
```

Birth of a program (C/Ux)



Building and running a program

```
chase:p0> make
```

```
gcc -l. -g3 -Wall -DNDEBUG -c dmm.c
```

```
...
```

```
1 warning generated.
```

```
gcc -l. -g -Wall -DNDEBUG -o basicdmmtest basicdmmtest.c dmm.o
```

```
gcc -l. -g -Wall -DNDEBUG -o test_basic test_basic.c dmm.o
```

```
gcc -l. -g -Wall -DNDEBUG -o test_coalesce test_coalesce.c dmm.o
```

```
gcc -l. -g -Wall -DNDEBUG -o test_stress1 test_stress1.c dmm.o
```

```
gcc -l. -g -Wall -DNDEBUG -o test_stress2 test_stress2.c dmm.o
```

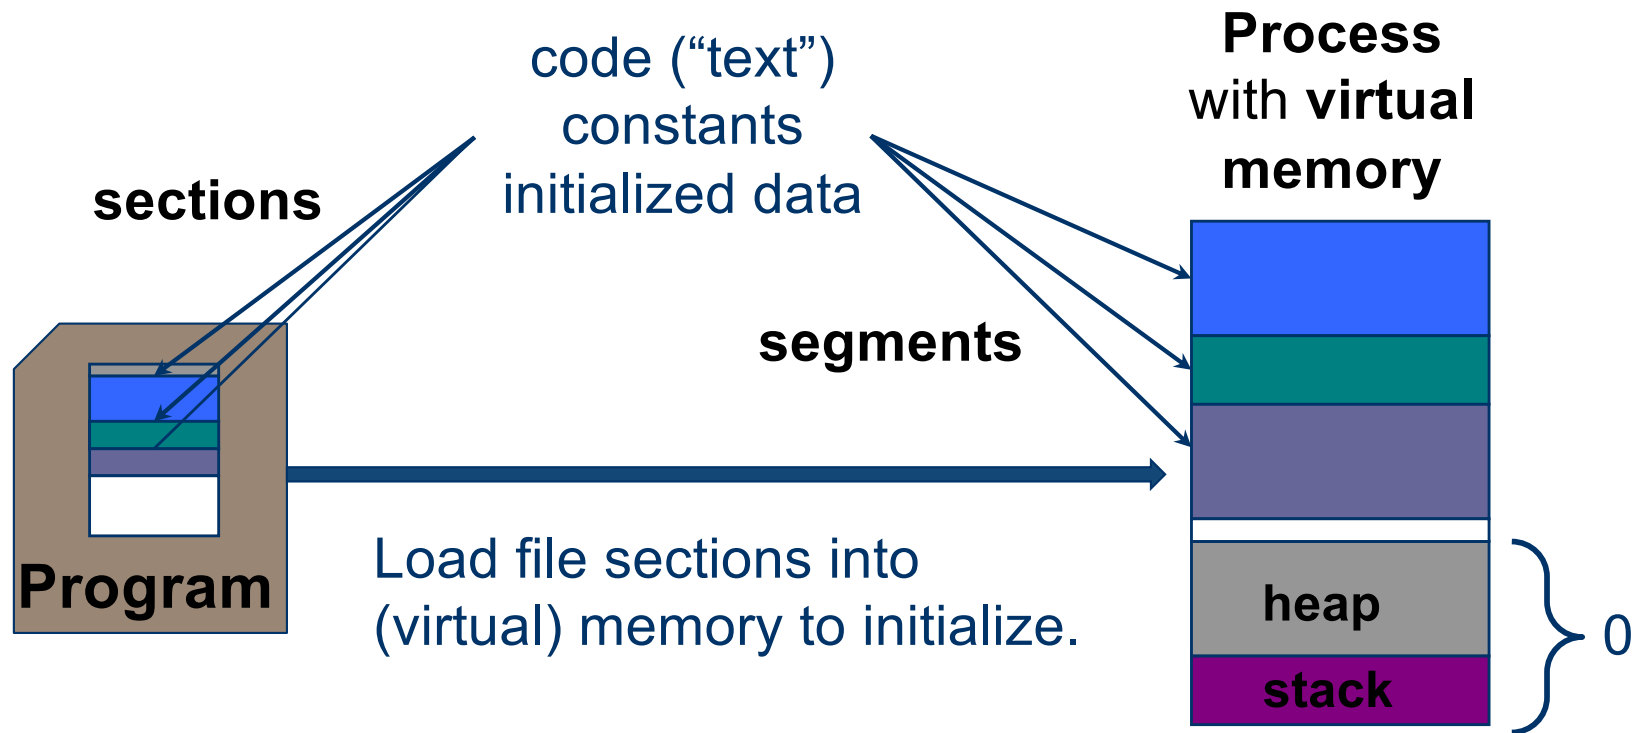
```
chase:p0> ./test_basic
```

```
calling malloc(10)
```

```
call to dmalloc() failed
```

```
chase:p0>
```

Running a program

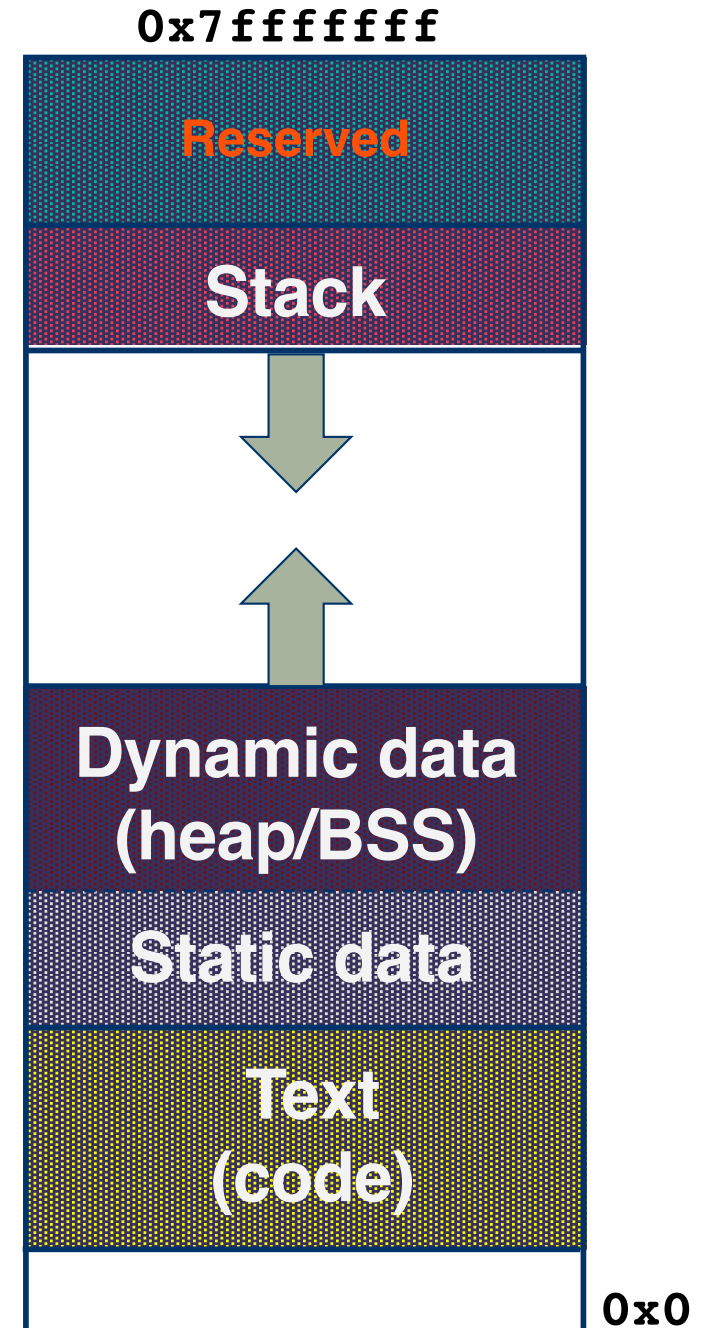


When a program launches, the OS initializes a **process** with a **virtual memory** to store the running program's code and data. Sections of the executable file initialize **segments** (regions) of the VM.

Process VM/VAS

(32-bit example)

- The program uses virtual memory through its process' **Virtual Address Space**:
- An addressable array of bytes...
- Containing every instruction it can execute...
- And every piece of data those instructions can reference...
 - E.g., read/write == **load/store** on memory
- Partitioned into logical **segments (regions)** with distinct purpose and use.
- Every reference by a running program is interpreted in the context of its VAS.
 - Resolves to a location in machine memory



VM segments: a view from C

- **Globals (static data):**
 - Fixed-size segment
 - Writable by user program
 - May have initial values
- **Text (instructions)**
 - Fixed-size segment
 - Executable
 - Not writable
- **Heap and stack**
 - Variable-size segments
 - Writable
 - Zero-filled on demand

