



D u k e S y s t e m s

**CPS 310**


**Threads in the Kernel: See How they Block**



**Jeff Chase**  
**Duke University**




# The story so far...



# Simplify for today

- For today, and for the thread labs, let's assume classic processes, with a single thread.
- Then we can use the terms “process” and “thread” interchangeably.
  - (As Linux people would do anyway.)
- Let's also assume a uniprocessor: one core, one slot.
  - Then we can ignore physical concurrency. There is only logical concurrency.
  - Which comes from....**what?**

# Two threads sharing a CPU/core



# The core-and-driver analogy



**zipcar**®  
wheels when you want them



**Core #1**



**Core #2**

The machine has a bank of CPU cores for threads to run on.

The OS allocates cores to threads (the “drivers”).

Cores are hardware. They go where the driver tells them.

OS can force a switch of drivers any time.

# What causes a context switch?

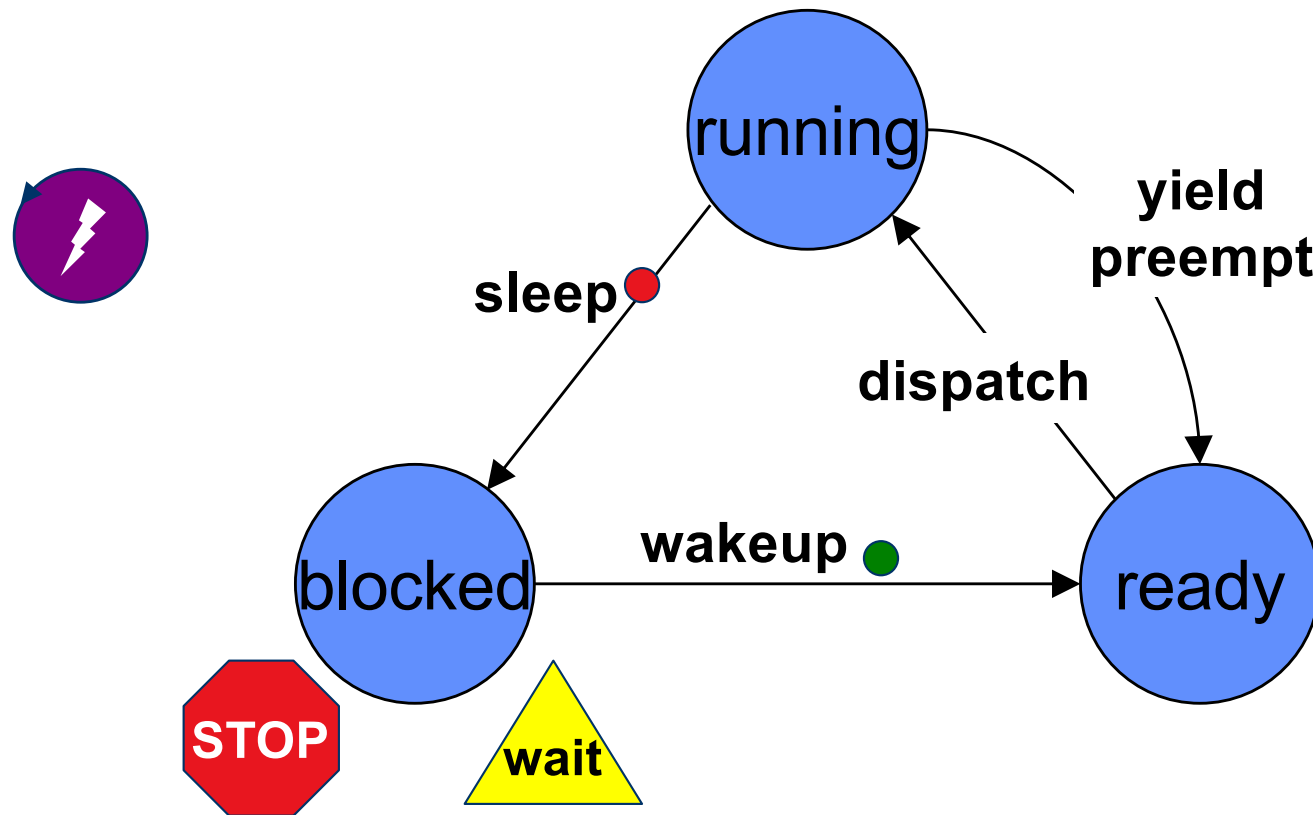
There are three possible causes:

1. Preempt (yield). The thread has had full use of the core for long enough. It has more to do, but it's time to let some other thread “drive the core”.
  - E.g., timer interrupt, quantum expired → OS forces yield
  - Thread enters Ready state, goes into pool of runnable threads.
2. Exit. Thread is finished: “park the core” and die.
3. Block/sleep/wait. The thread cannot make forward progress until some specific occurrence takes place.
  - Thread enters Blocked state, and just lies there until the event occurs. (Think “stop sign” or “red light”.)



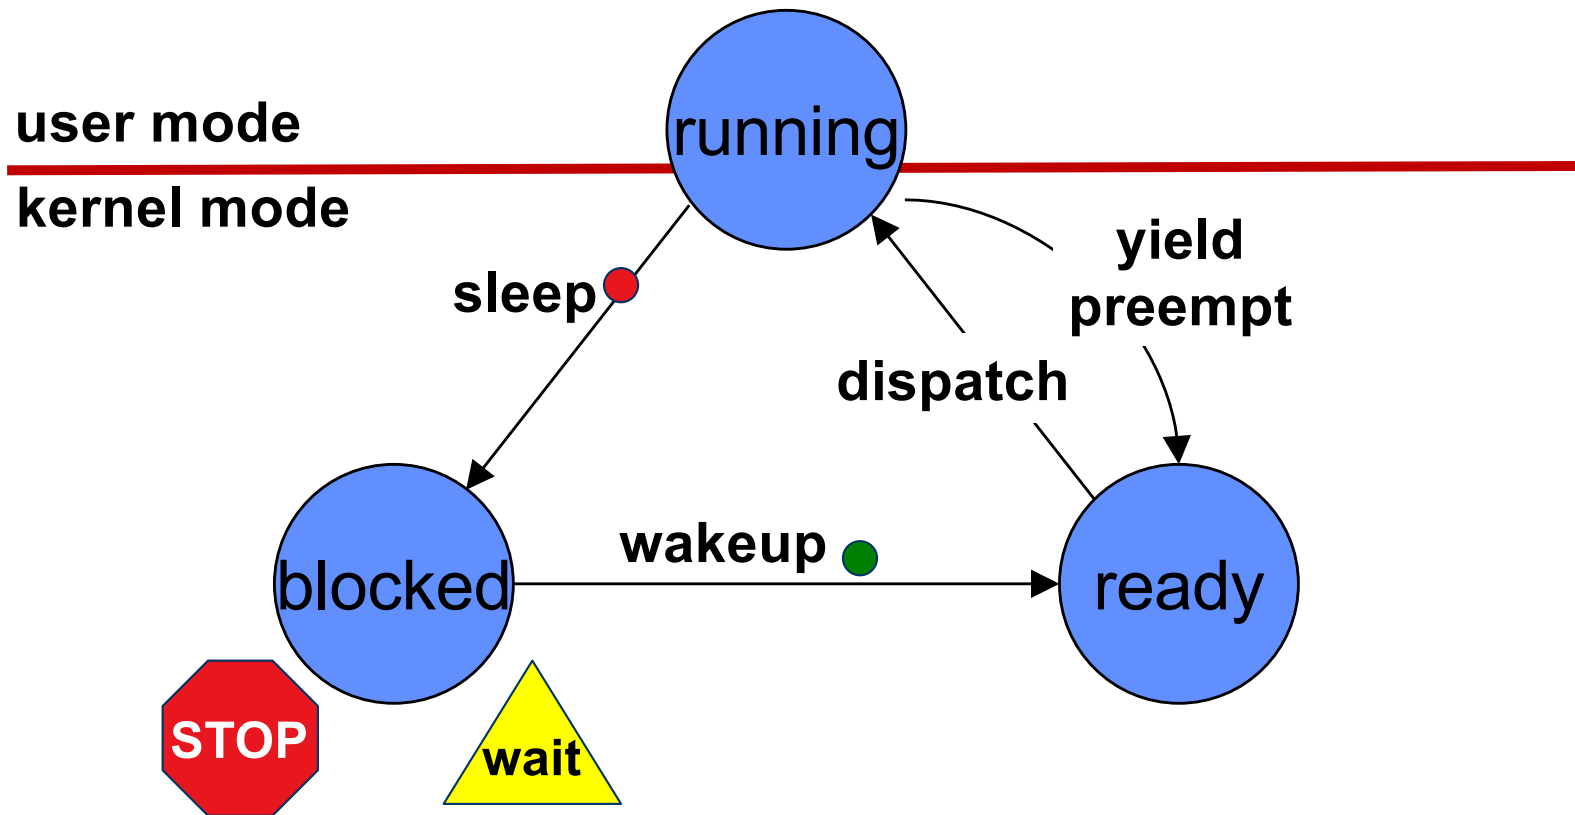
# Thread states

Concept: block/sleep/wait



Often a thread's execution is blocked: it cannot continue until a designated event occurs. The thread waits for the event. We say that the thread **blocks** or **sleeps** and that the event **wakes** it up.

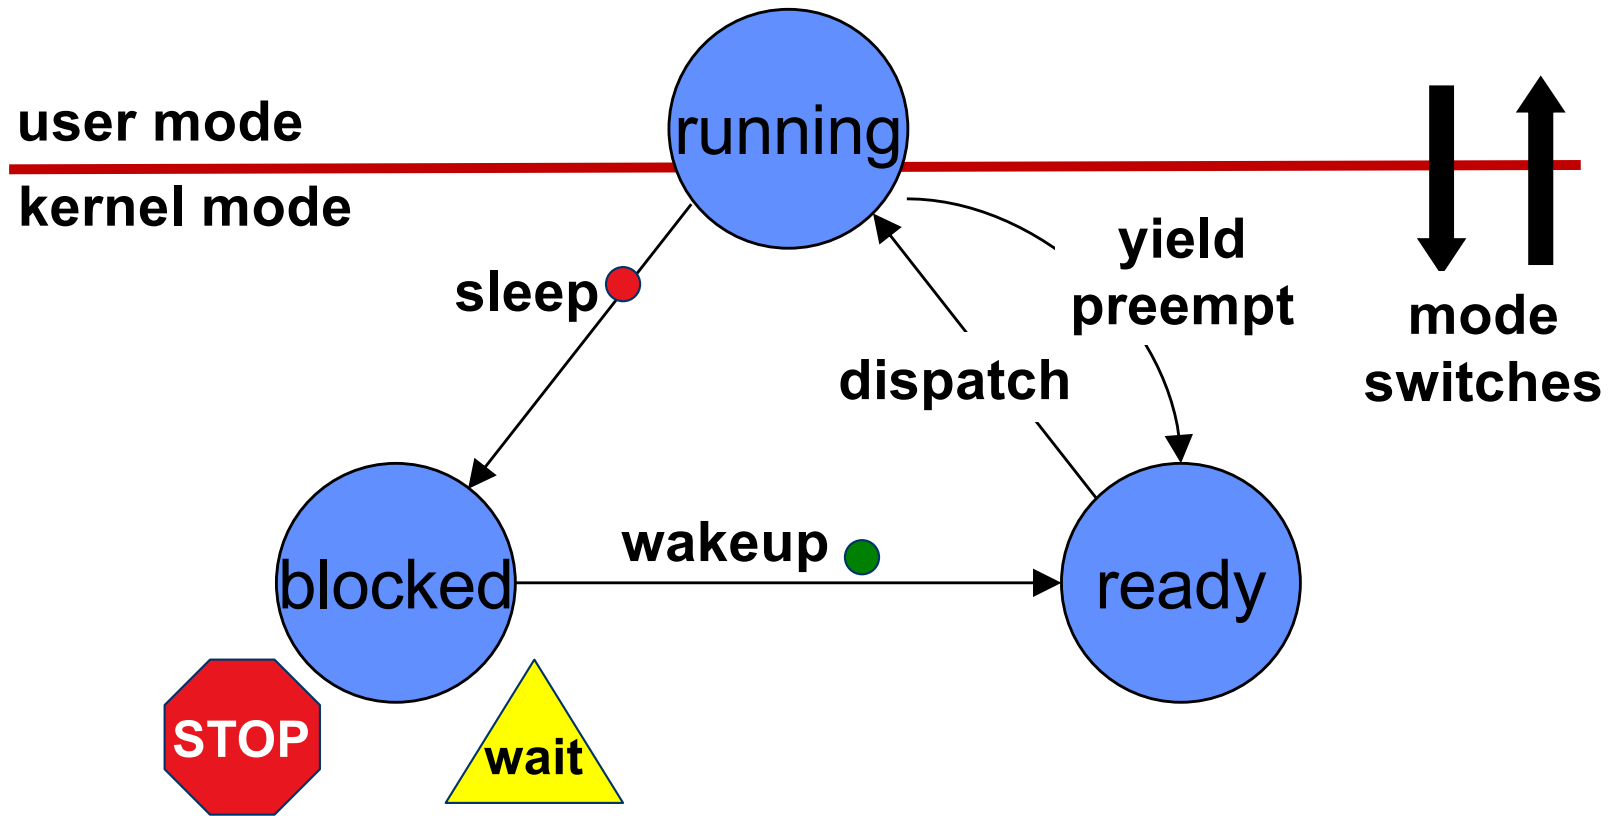
# Thread states: user mode vs. kernel



In a “real” OS, threads/processes change states only in kernel mode.  
Let us assume for today that our threads are in kernel mode



# Mode switches

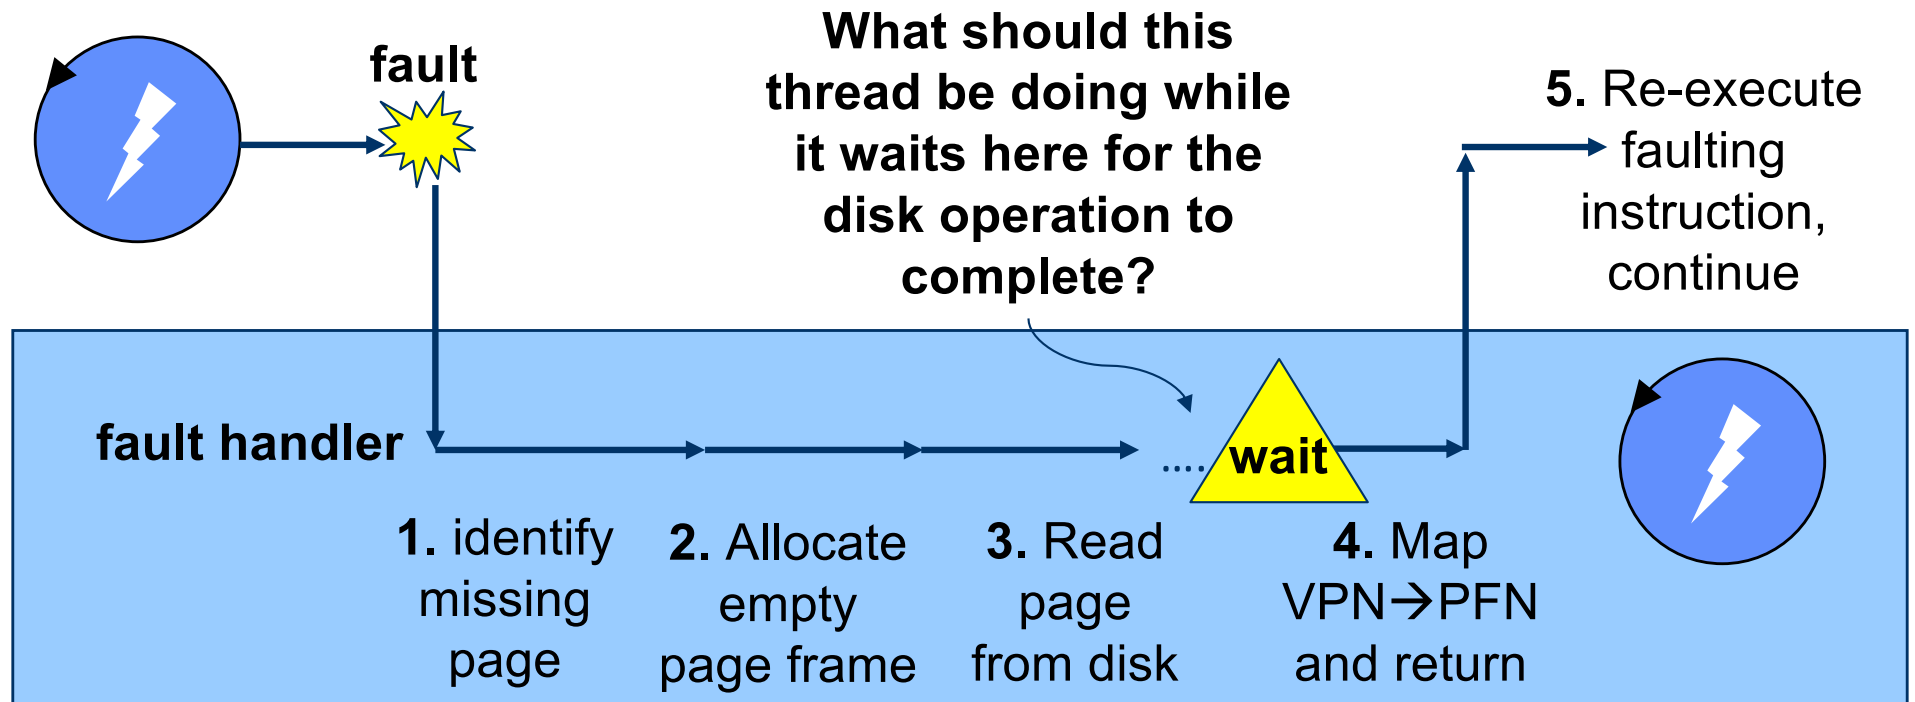


## Mode switches


User mode to kernel mode: trap, fault, or interrupt.

Kernel mode to user mode: return from trap, fault, or interrupt.

# Handling a page fault



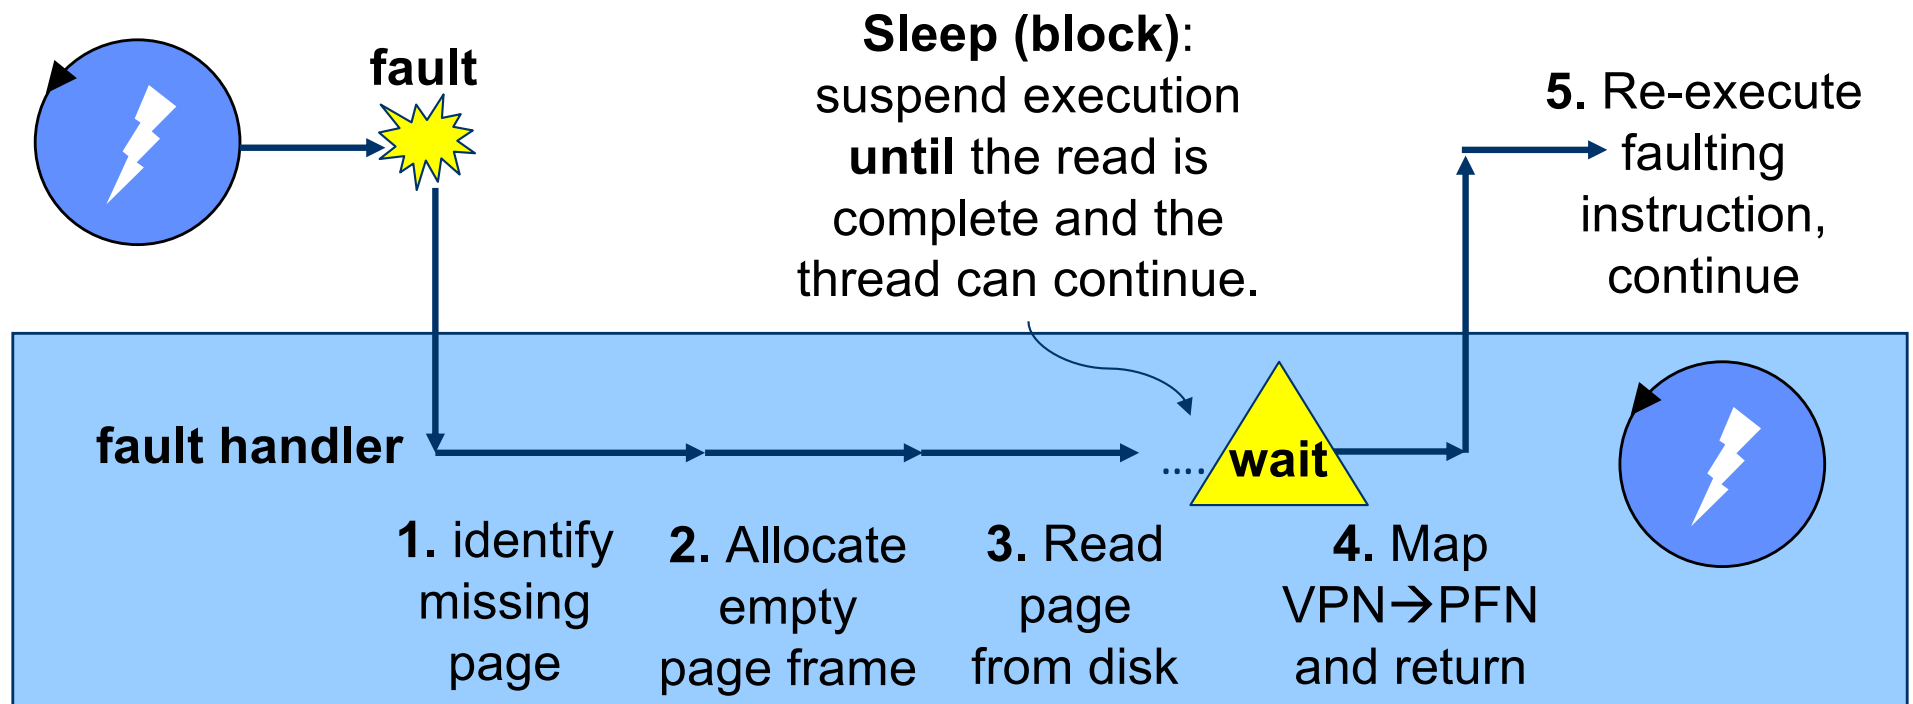
# Blocking/sleeping in a page fault



When a thread **blocks**, it stops running, leaves the core for use by other threads, and **sleeps**: it wakes/resumes only after some specified event or condition occurs.

Threads block for page faults, and for many other reasons as well.  
**It happens all the time.**

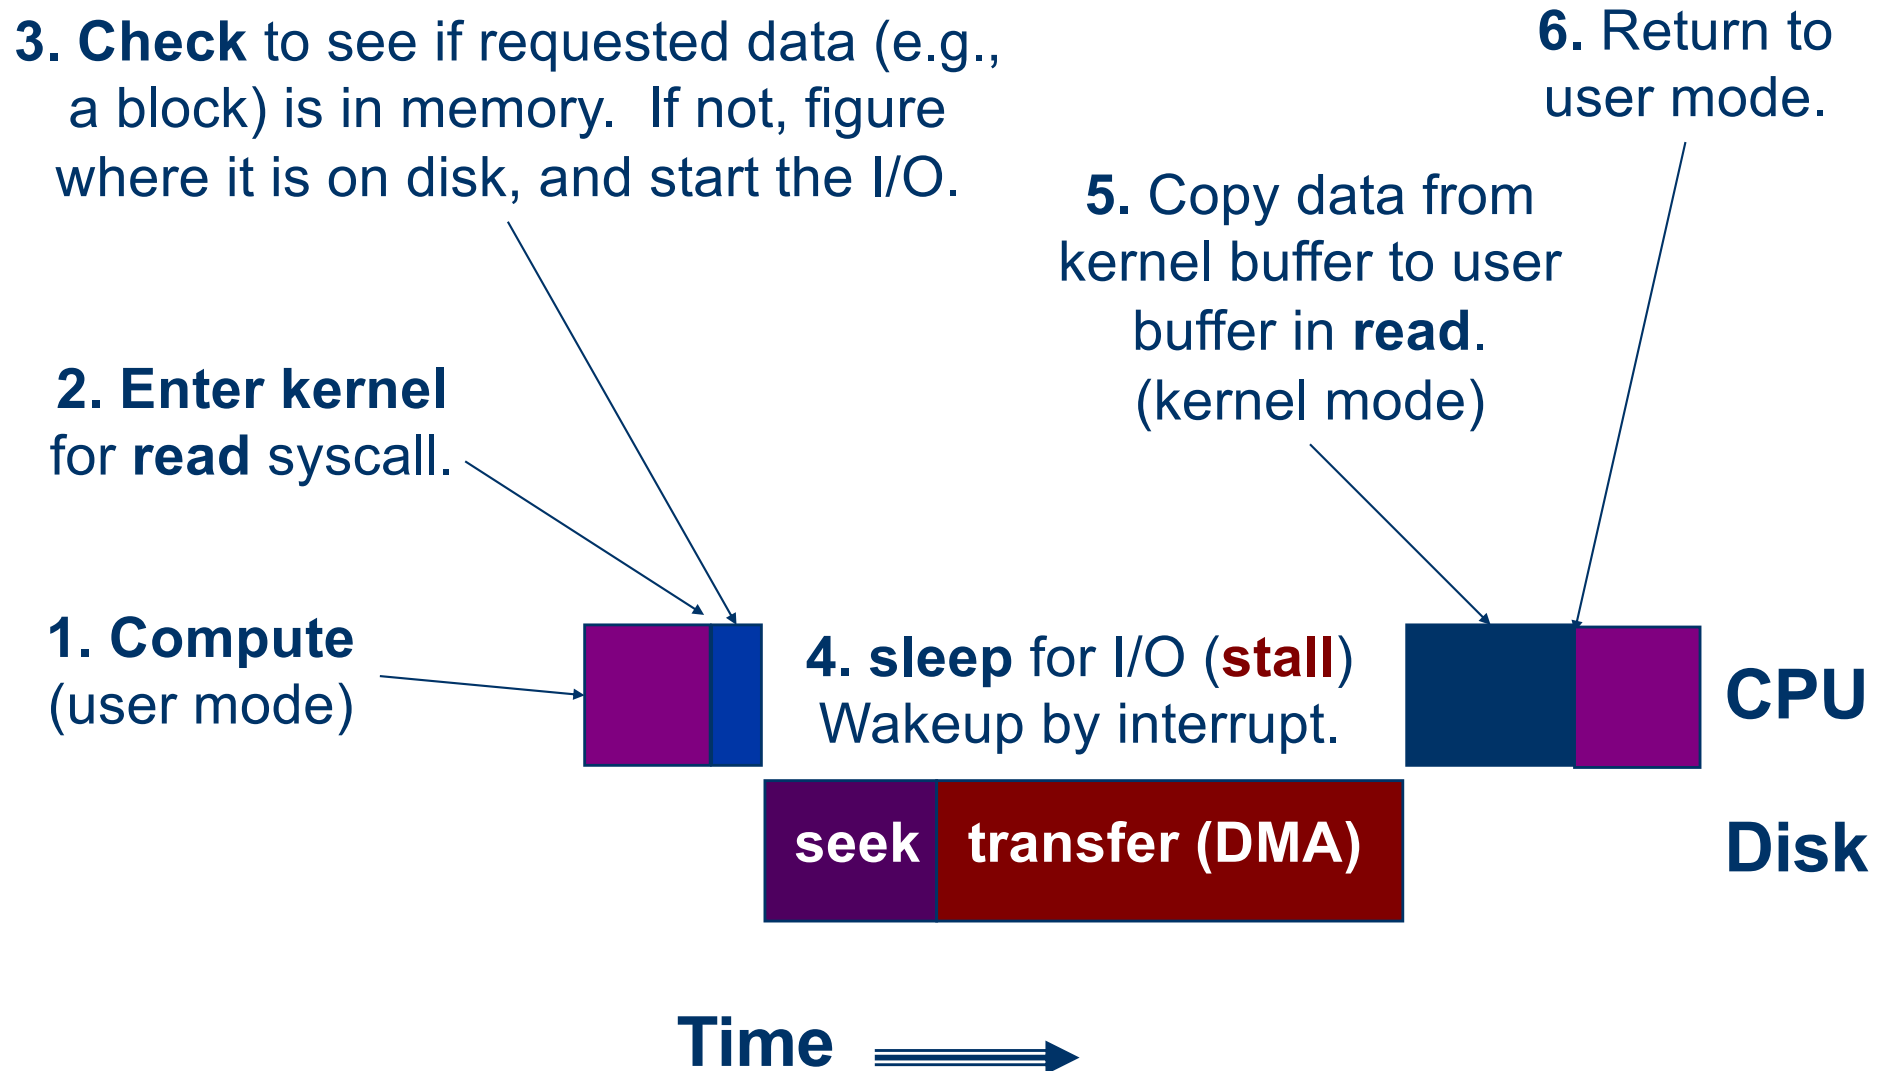
# Blocking/sleeping in a page fault



When a thread **blocks**, it stops running, leaves the core for use by other threads, and **sleeps**: it wakes/resumes only after some specified event or condition occurs.

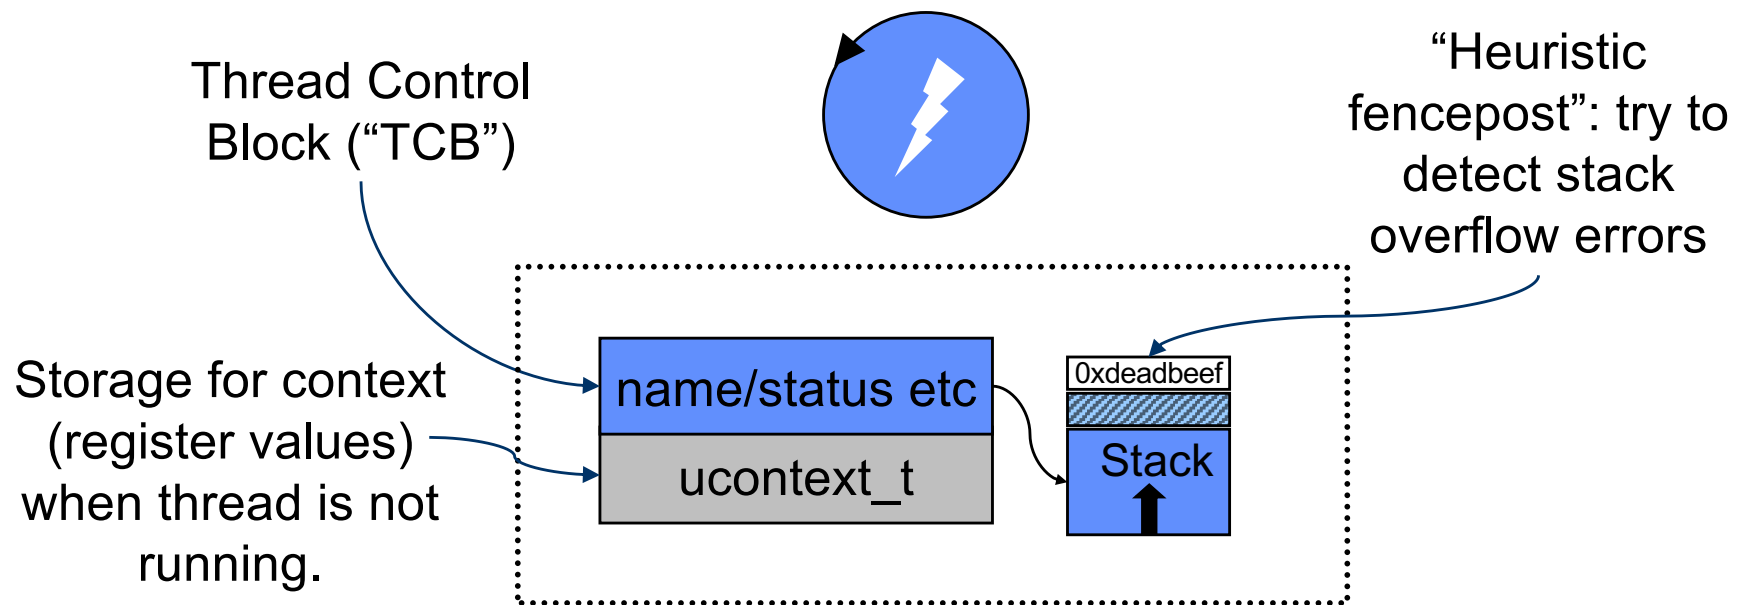
Threads block for page faults, and for many other reasons as well.  
**It happens all the time.**

# Anatomy of a read()




# Portrait of a thread: the TCB

In an implementation, each thread is represented by a data structure. We call it a “**thread object**” or “**Thread Control Block**”. It stores information about the thread, and may be linked into other system data structures.



Each thread also has a runtime stack for its own use. As a running thread calls procedures in the code, frames are pushed on its stack.

# Kernel Stacks and Trap/Fault Handling





The syscall (trap) handler makes an indirect call through the system call dispatch table to the handler registered for the specific system call.

# More analogies: threads and stacks

- Each thread chooses its own path. (By its program.)
- But they must leave some “bread crumbs” to find their way back on the return!
- Where does a thread leave its crumbs? **On its stack!**
  - Call frames with local variables
  - Return addresses

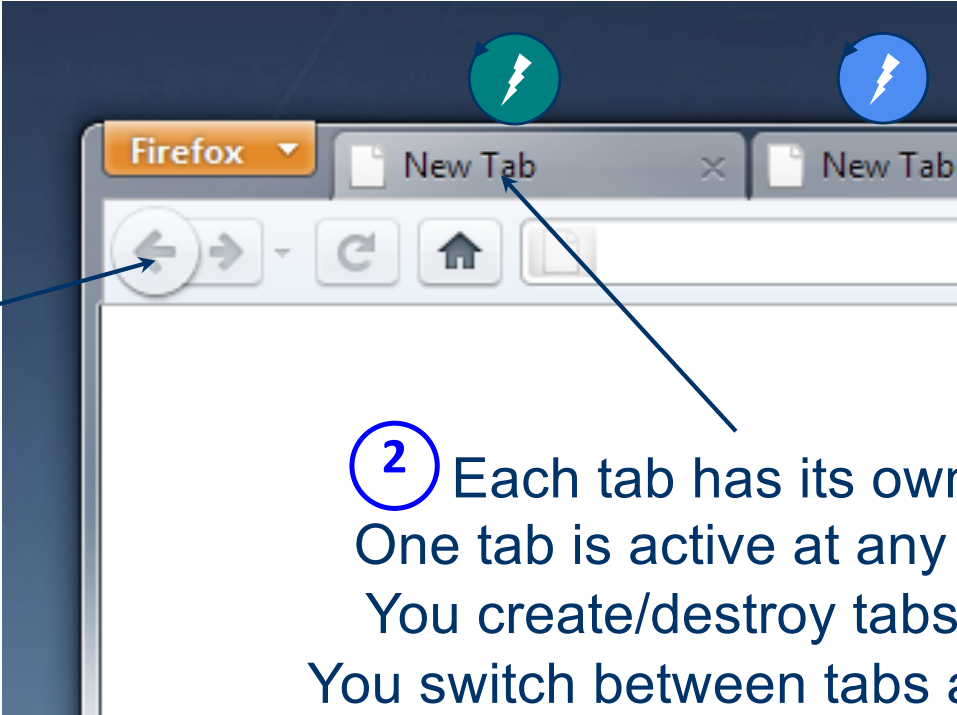
This means that **each thread must have its own stack** in memory, so that their crumbs aren't all mixed together.





# More analogies: context/switching


① Page links and back button navigate a “stack” of pages in each tab.



② Each tab has its own stack. One tab is active at any given time. You create/destroy tabs as needed. You switch between tabs at your whim.



- ③ Similarly, each thread has a separate stack. The OS switches between threads at its whim. One thread is active per CPU core at any given time.




# Thread states and queues

- At most on thread T is **running** (per core).
  - Current[core] points at T's TCB.
- Any other thread P is **blocked** or **ready**.
- Link P's TCB onto a **thread queue**.

Internal routines manage thread state changes.  
E.g.: **yield**, **sleep**, **wakeup**.

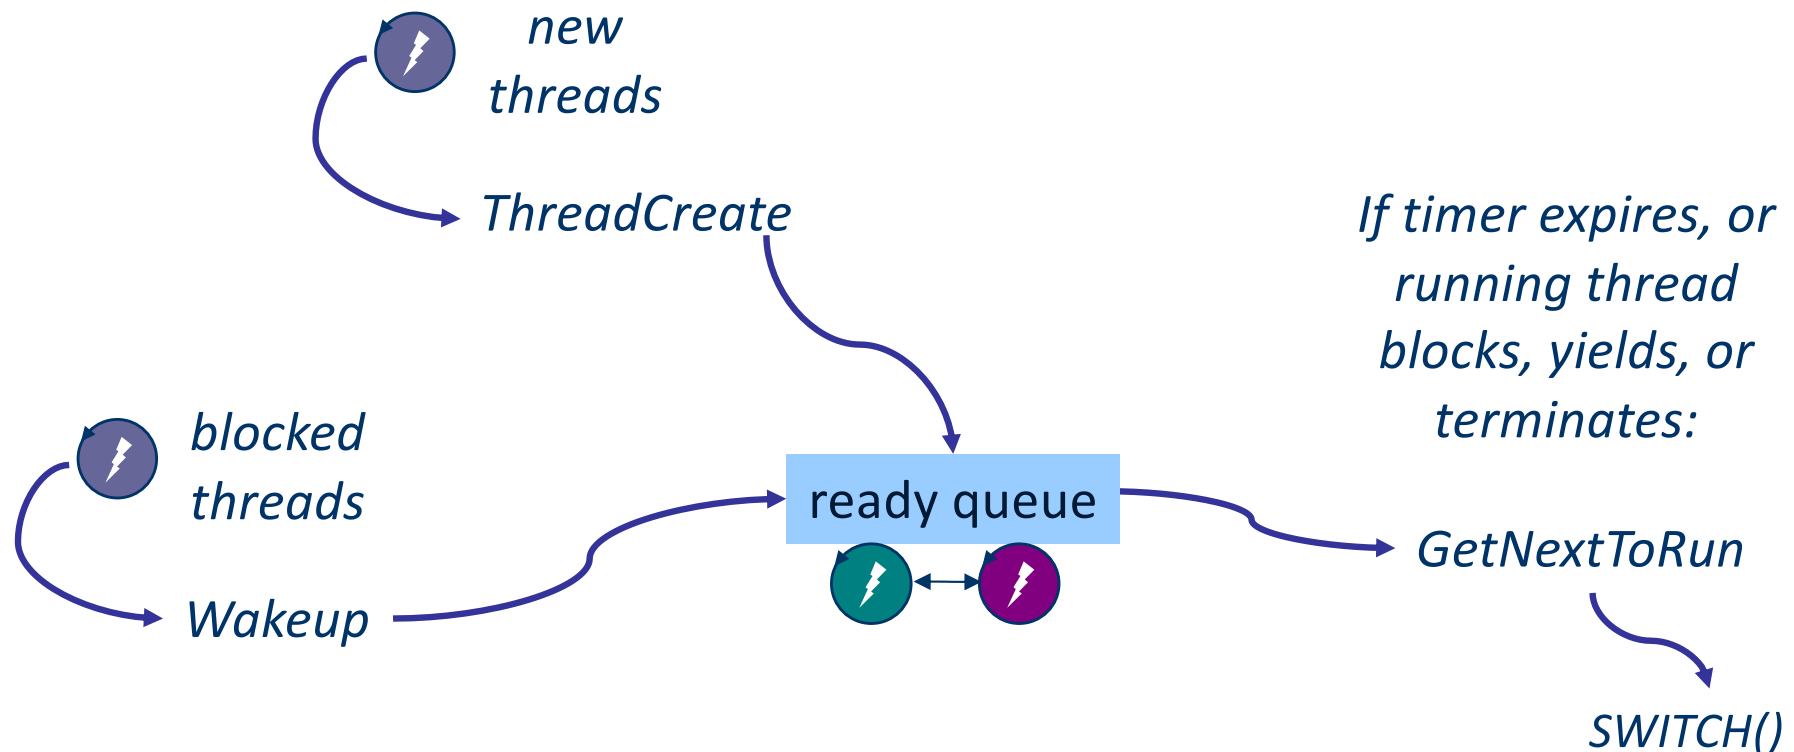
These routines slam thread TCBs on and off various internal thread queues.



**Ready** → ready queue

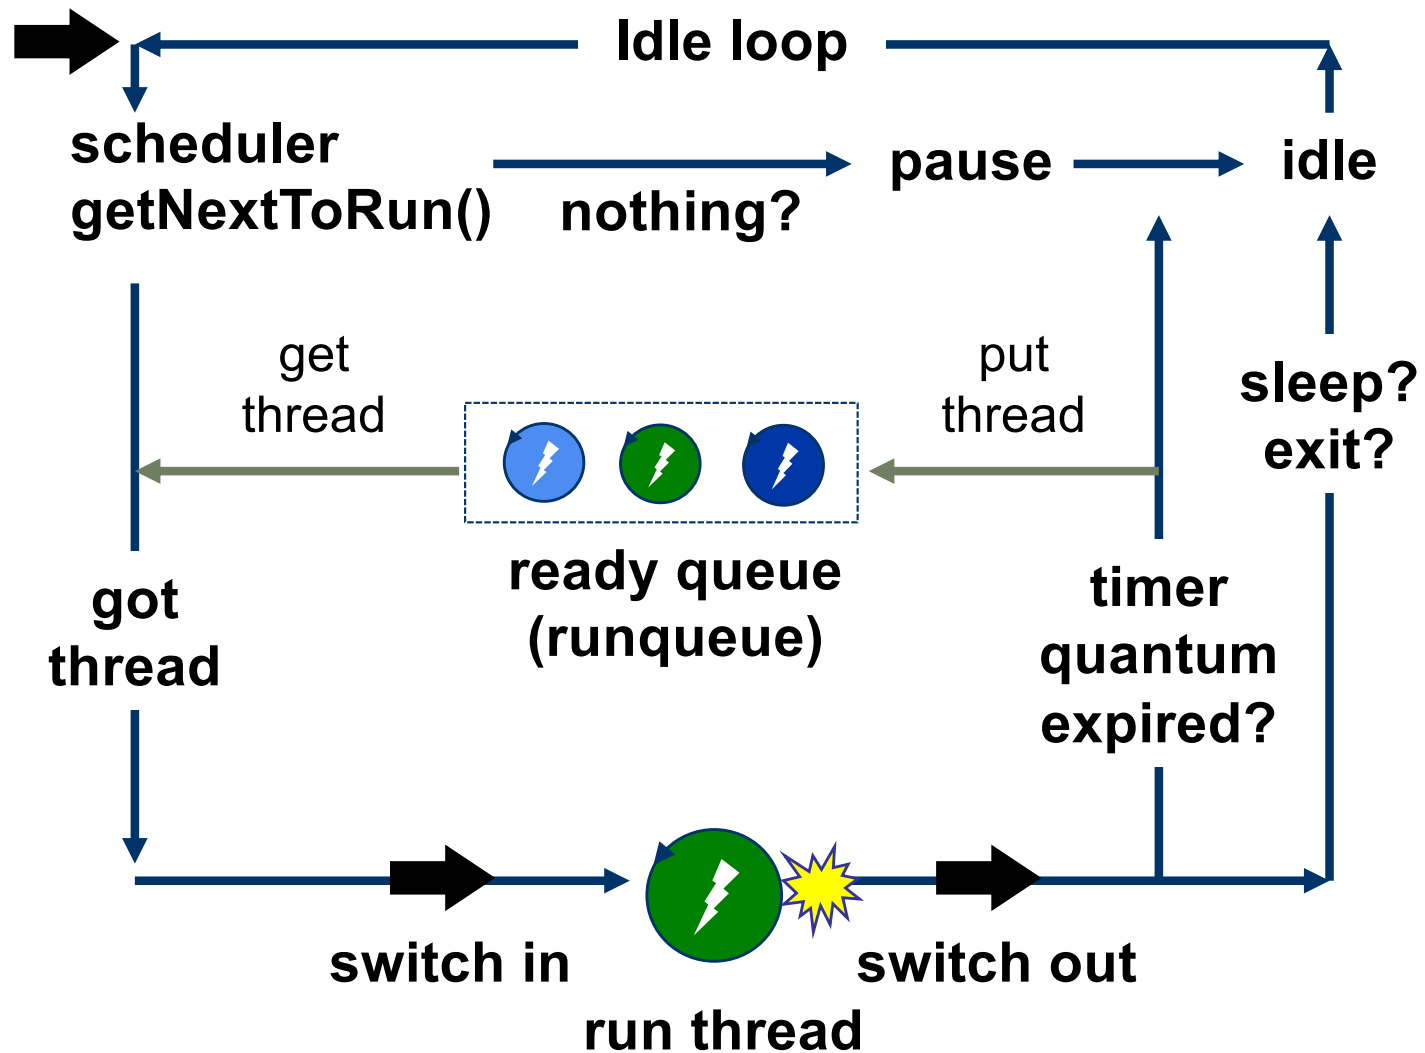
**Blocked** → waiter queue for event/object


# The ready queue



For p1t, the ready pool is a simple FIFO queue: the **ready list** or **ready queue (runqueue)**. Scalable multi-core systems use multiple runqueues to reduce locking contention among cores. It is typical to implement a ready pool as an array of runqueues for different priority levels.

# What cores do





A former student named Grace Chen drew this artwork freeform for extra credit on a final exam.